Privacy-preserving Machine Learning

Machine Learning

Image processing

Speech recognition

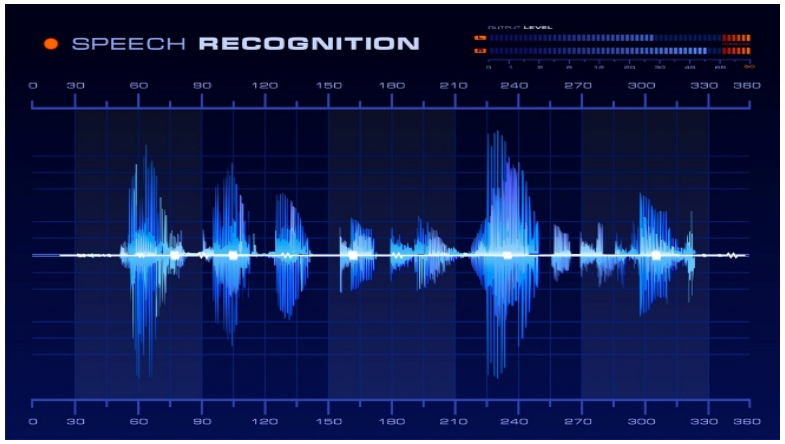

Playing Go

OpenAI

More data → Better Models

Map Predictions

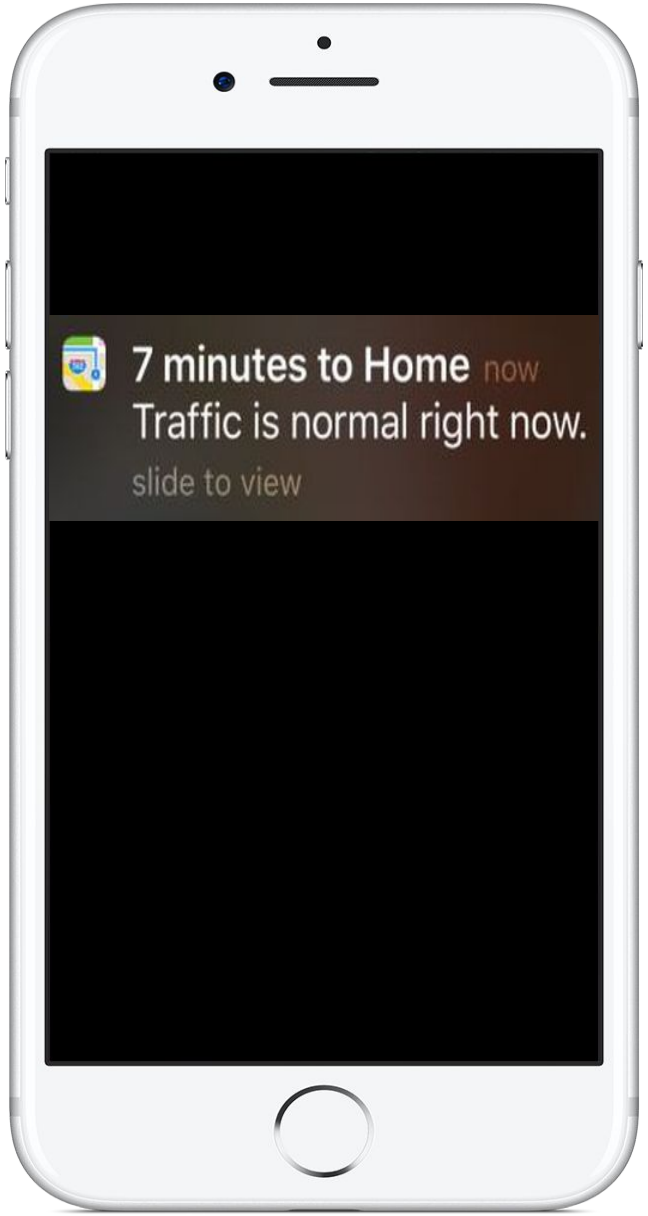

Customized Homepage

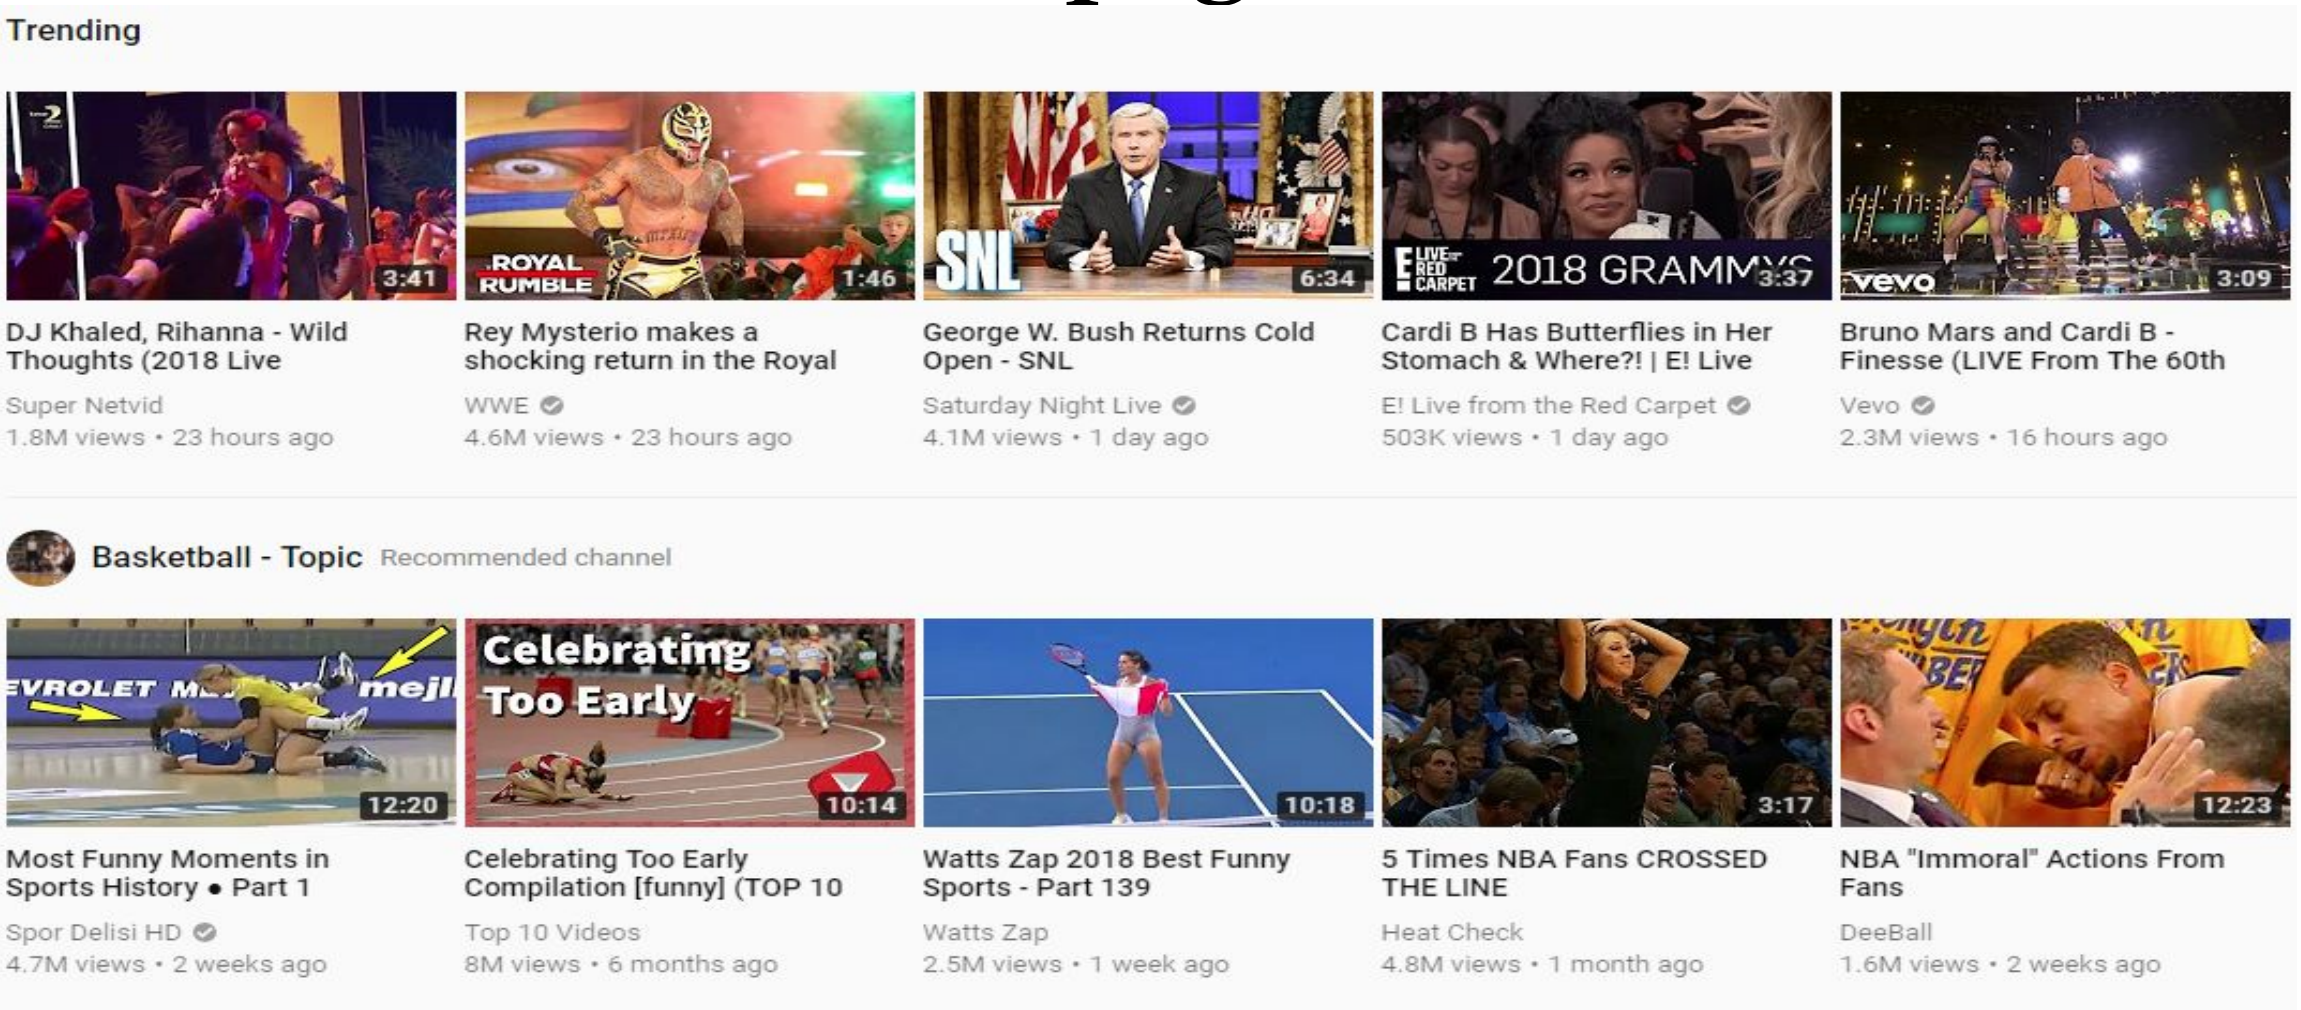

✔ Nice machine learning applications benefiting our lives

× Models trained on sensitive data

Can companies train the models without learning our data 

Privacy-preserving Machine Learning

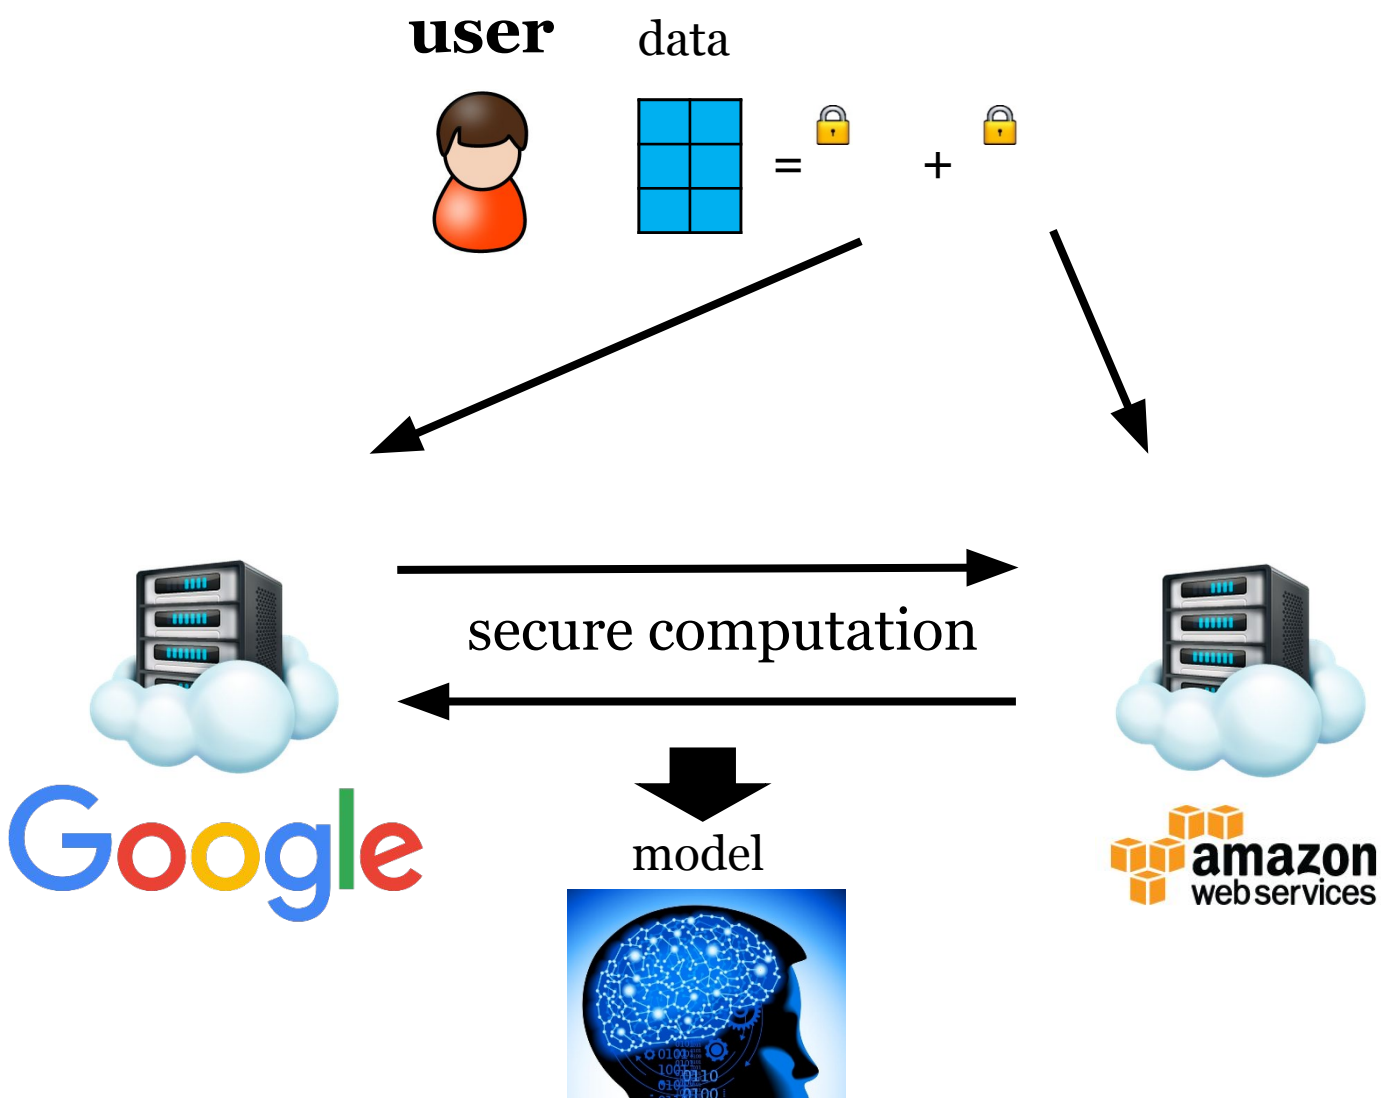

Use Case for Companies: Fraud detection

Card #	Time	Location	Amount
xxxxxxx	8/8/2016	CA, USA	xx.xx
.....			

Name	SSN
Alice	xxxxx
.....	

xxxxxxx	x/xx/xxxx	xx,xxx	xx.xx
.....			

Alice	xxxxx
.....	

Use Case for Hospitals

ID	Name	Sex	DOB	Disease xxx
101	Alice	F	1976-2-23	N
102	Bob	M	1992-10-12	N
103	Charlie	M	1983-1-15	N
104	David	M	2005-4-30	Y
…	…	…	…	N

In theory, there is no difference between theory and practice;

In practice, there is.

Linear Regression

Linear Regression

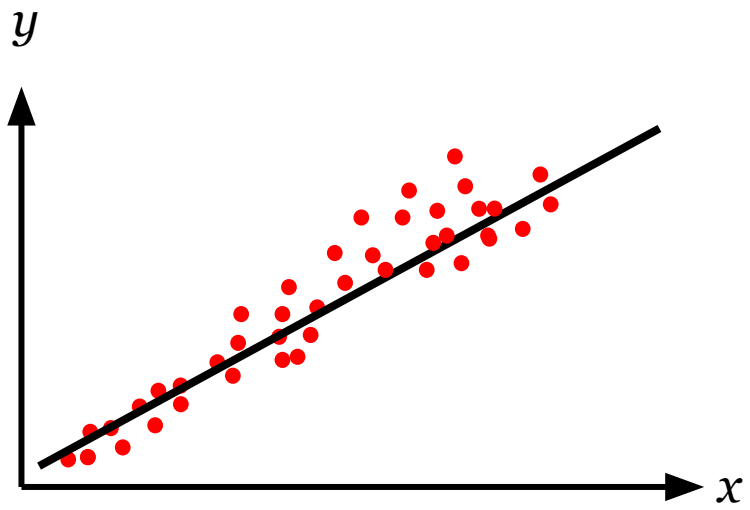

Input: data value pairs (x, y)s

Output: model w

$$y^* = \sum_i w_i x_i = w \cdot x \approx y$$

Cost function (Loss function)

$$y^* = \sum_i w_i x_i = w \cdot x \approx y$$

- $$C_x(w) = \frac{1}{2}(y^* - y)^2$$

$$C(w) = \frac{1}{n}\sum_x C_x(w)$$

$$\arg\min_w C(w)$$

Closed-form solution for linear regression

- $$w = (X^T X)^{-1} X^T y$$

O(n^3), slow for large datasets

Gradient decent

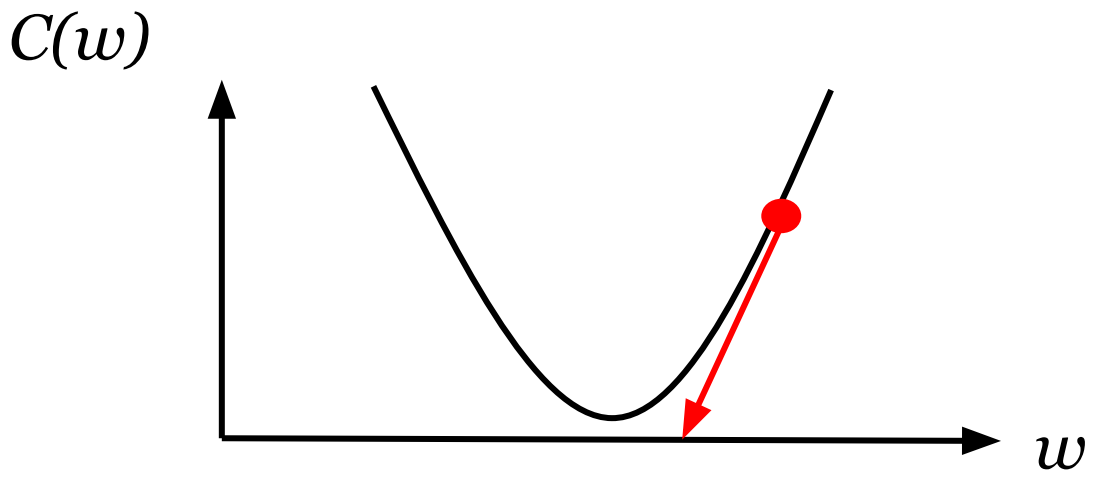

$$y^* = \sum_i w_i x_i = w \cdot x$$

$$C_x(w) = \frac{1}{2}(y^* - y)^2$$

$$C(w) = \frac{1}{n}\sum_x C_x(w)$$

1. Initialize w randomly
2. Compute derivative of $C(w)$
3. Update w

$$w = w - \alpha \frac{1}{n}\sum_x (x \cdot w - y)x \qquad w_i = w_i - \alpha \frac{1}{n}\sum_x (x \cdot w - y)x_i$$

Stochastic gradient decent (SGD)

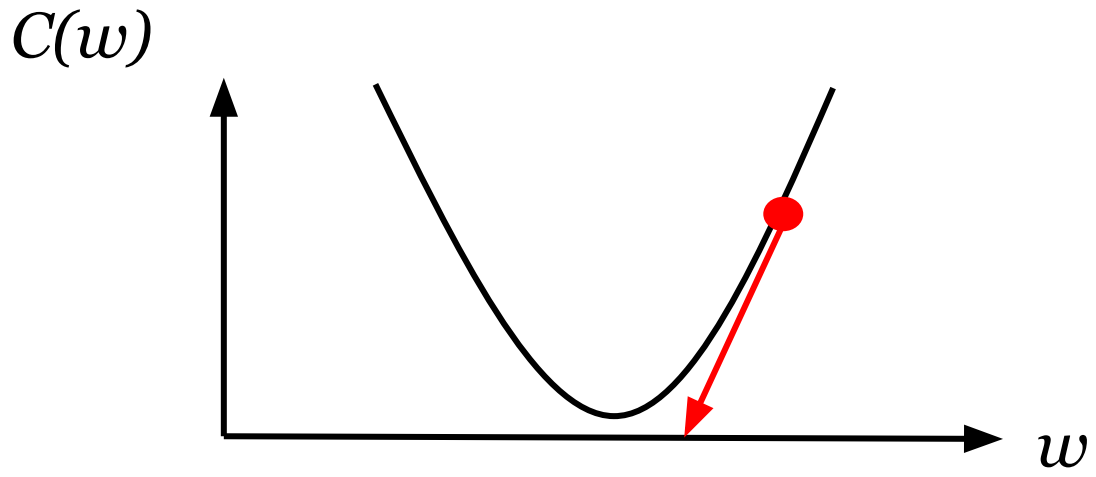

1. Initialize w randomly
2. Select a random sample (x, y), compute derivative of $C_x(w)$
3. Update w

$$w = w - \alpha(x \cdot w - y)x$$

$$w_i = w_i - \alpha(x \cdot w - y)x_i$$

Mini-batch SGD

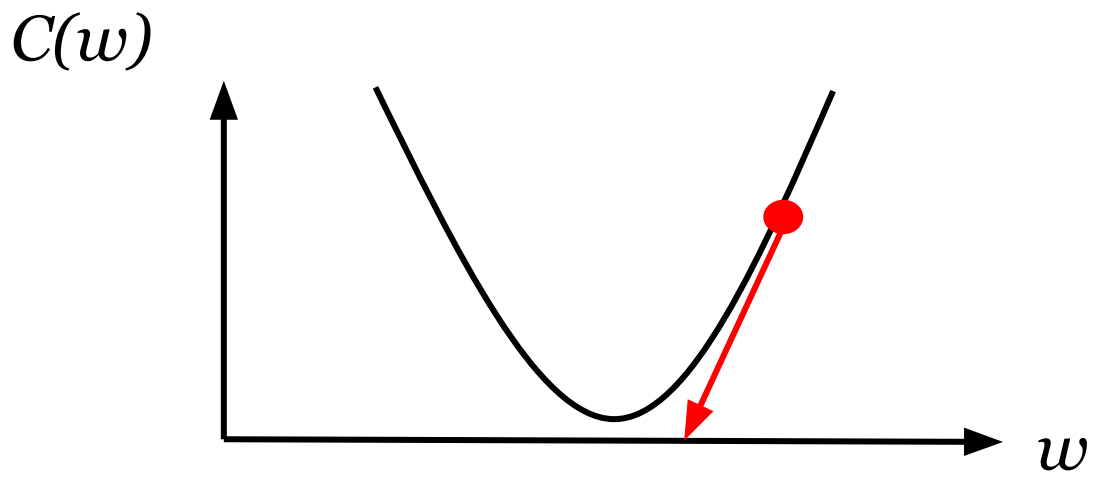

1. Initialize w randomly
2. Select a batch of random samples (x, y), compute derivative
3. Update w

$$w = w - \alpha \frac{1}{|B|} \sum_{x \in B} (x \cdot w - y) x$$

Mini-batch SGD

- Epoch: randomly shuffle all the data, select |B| samples each round
- Vectorization:

$$w = w - \alpha \frac{1}{|B|} \sum_{x \in B} (x \cdot w - y) x$$

$$w = w - \frac{\alpha}{|B|} \cdot X^{\mathrm{T}} \times (X \times w - y)$$

Other variants

$$y^* = \sum_i w_i x_i + b = \mathbf{w} \cdot \mathbf{x} + b$$

Ridge regression: $C_x(w) = \frac{1}{2}(y^* - y)^2 + \lambda \left|\left|w\right|\right|_2^2$

Adaptive learning rate α

Privacy-preserving linear regression

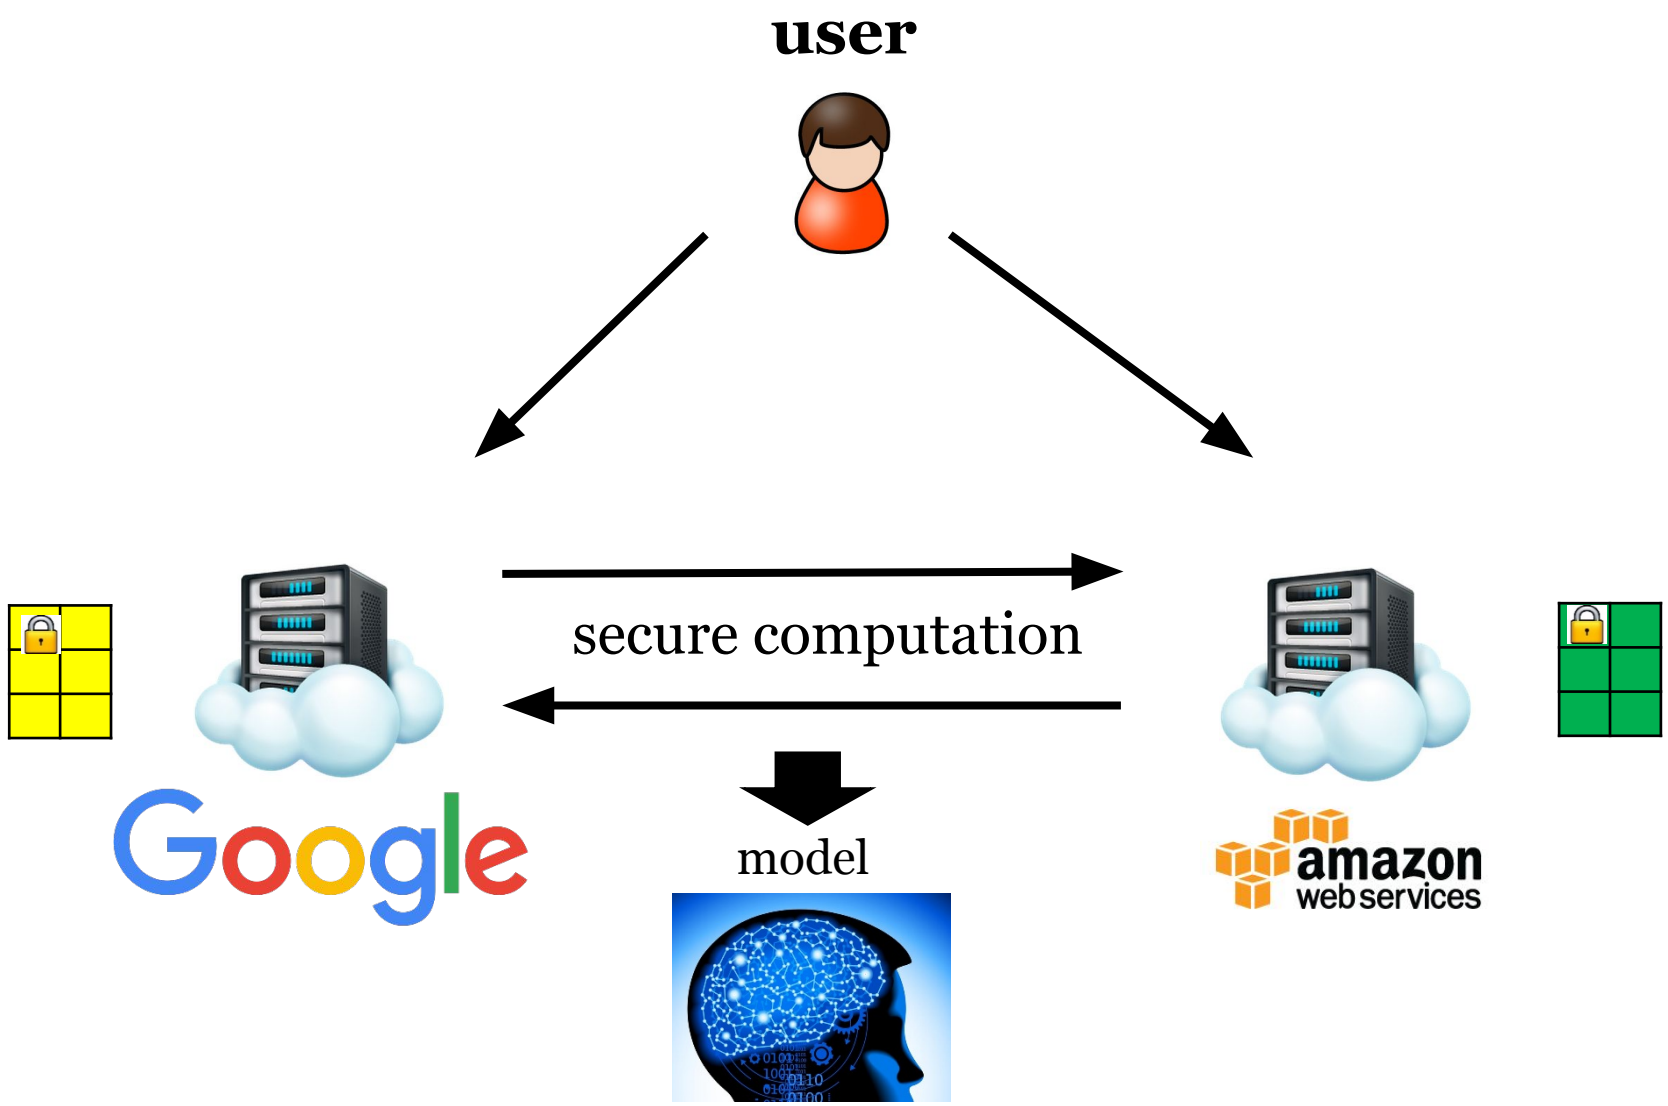

Privacy-preserving linear regression

SGD: $w_i = w_i - \alpha(x \cdot w - y)x_i$

1. Users secret share data and values (x,y)
2. Servers initialize and secret share the model w
3. Run SGD using GMW protocol

Decimal number?

Fixed-point multiplication

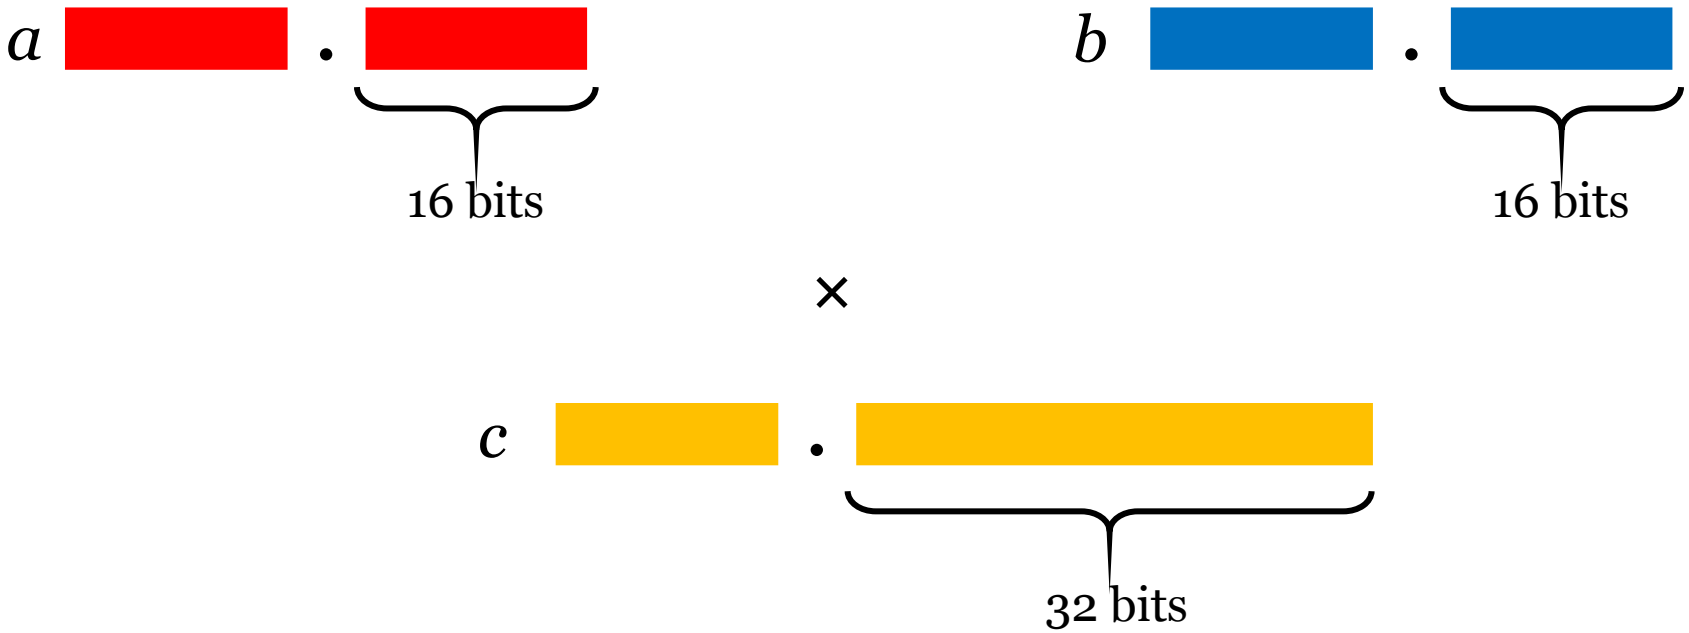

- Truncation: Same as integer multiplication
- Decimal part grows → overflow

16 bits

fixed-point multiplication

Truncation on shared values

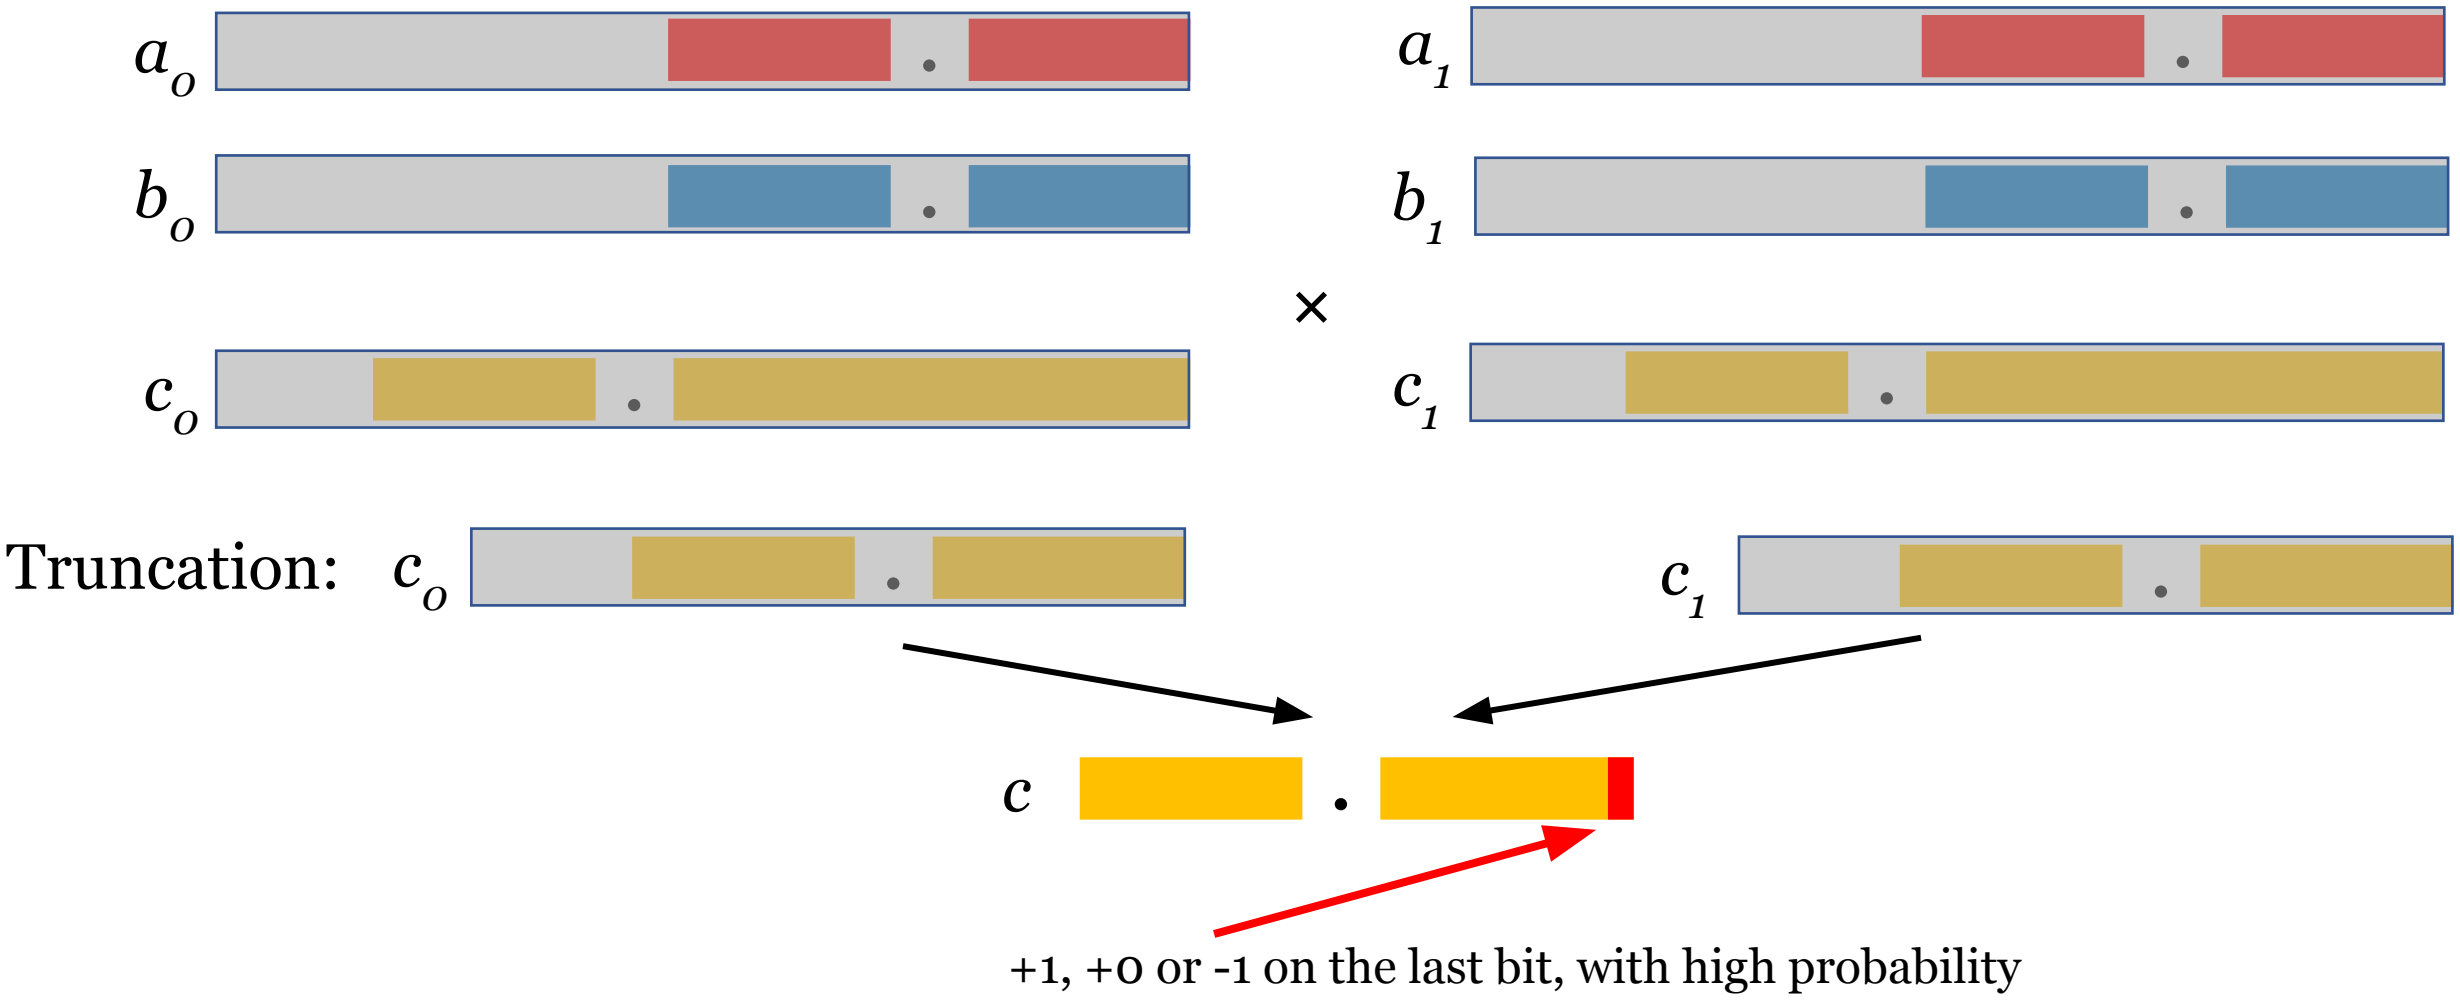

Privacy-preserving linear regression

SGD: $w_i = w_i - \alpha(x \cdot w - y)x_i$

1. Users secret share data and values (x,y)
2. Servers initialize and secret share the model w
3. Run SGD using GMW protocol
4. Truncate the shares after every multiplication

Effects of truncation

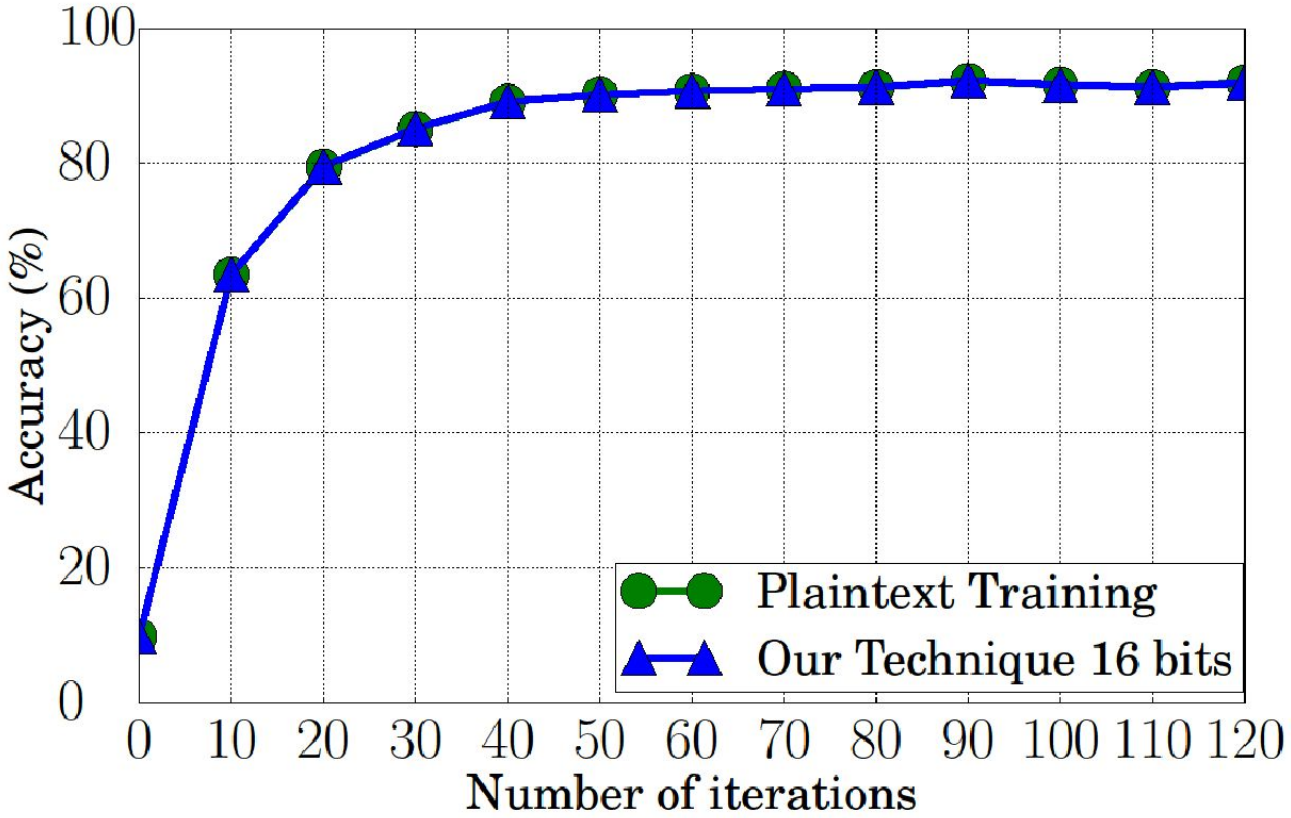

- 4-8× faster than fix-point multiplication garbled circuit